

# 06 Implement SGD

April 19, 2019

## 0.1 Mini-Batch SGD implementation on Boston Housing dataset

### [1.1] Import Statements

```
In [868]: import warnings
          warnings.filterwarnings("ignore")
          from sklearn.datasets import load_boston
          from random import seed
          from random import randrange
          from csv import reader
          from math import sqrt
          from math import pow
          from sklearn import preprocessing
          import pandas as pd
          import numpy as np
          import matplotlib.pyplot as plt
          from prettytable import PrettyTable
          from sklearn.linear_model import SGDRegressor
          from sklearn import preprocessing
          from sklearn.metrics import mean_squared_error
          import random
          from sklearn.preprocessing import StandardScaler
          from sklearn.model_selection import train_test_split
```

### [1.2] Load Dataset

```
In [869]: X = load_boston().data
          Y = load_boston().target
```

### [2.1] Train Test Split

```
In [870]: X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.3, random_state=42)
          scaler = preprocessing.StandardScaler().fit(X_train)
          X_train = scaler.transform(X_train)
          X_test = scaler.transform(X_test)
```

### [3.1] Applying Mini-Batch SGD

```
In [871]: def plott(Y_samples, y_pred, title):
          plt.scatter(Y_samples, y_pred)
          plt.xlabel('Y sample')
          plt.ylabel('Y pred')
          plt.title(title)
```

```
In [872]: #Source: https://stackoverflow.com/questions/15923826/random-row-selection-in-pandas
          # https://stackoverflow.com/questions/17260109/sample-two-pandas-dataframes-the-same
```

```
def random_sampling(X, Y, n_samples):
    X = pd.DataFrame(X)
    X_samples = X.sample(n=n_samples)
    Y_samples = Y[X_samples.index]
    Y_samples = pd.DataFrame(Y_samples)
    return X_samples, Y_samples
```

```
In [873]: #Source: https://machinelearningmastery.com/implement-linear-regression-stochastic-gradient-descent-from-scratch-in-python/
          # https://stackoverflow.com/questions/50328545/stochastic-gradient-descent-for-linear-regression
```

```
def stochastic_gradient_descent(X_data, y_data, epochs, n_samples, learning_rate_given):
    w_0 = np.random.normal(0,1, 13)
    eta0 = 0.01
    power_t = 0.5
    t = 1
    b_0 = np.random.rand()
    for epoch in range(epochs):
        if epoch == 0:
            w_i = w_0
            b_i = b_0
        else:
            w_i = w_latest
            b_i = b_latest

        #sampling data
        X_samples, Y_samples = random_sampling(X_data, y_data, n_samples)
        X_samples = X_samples.as_matrix()
        Y_samples = Y_samples.as_matrix()

        if learning_rate_given == 'constant':
            eta = eta0
        elif learning_rate_given == 'invscaling':
            if epoch == 0:
                eta = eta0
            else:
                t += 1
                eta = eta0 / (t**power_t)
```

```

        # Initializing derivatives to zero
        m_derivative = np.zeros(13)
        b_derivative = 0

        for X_sample, Y_sample in zip(X_samples, Y_samples):
            # m_derivative = -2*x_i(y_i - (w_T*x_i + b_i))
            # b_derivative = -2*(y_i - (w_T*x_i + b_i))
            y_hat = np.dot(w_i.T, X_sample) + b_i
            error = Y_sample - y_hat
            m_derivative += -(2/n_samples) * X_sample * error
            b_derivative += -(2/n_samples) * error

            w_latest = m_derivative * eta
            b_latest = b_derivative * eta
            w_latest = w_i - w_latest
            b_latest = b_i - b_latest
        y_pred = []
        for X_sample in X_samples:
            y_pred.append(np.dot(w_latest.T, X_sample) + b_latest)
        train_error = mean_squared_error(Y_samples, np.array(y_pred))
        return X_samples, Y_samples, y_pred, w_latest, b_latest, train_error

```

In [874]: `def mean_sq_error(X_test, y_test, w_latest, b_latest, plot_title):`

```

    X_test = pd.DataFrame(X_test)
    y_test = pd.DataFrame(y_test)

    X_test = X_test.as_matrix()
    y_test = y_test.as_matrix()

    y_pred = []
    for item in X_test:
        y_pred.append(np.dot(w_latest.T, item)+b_latest)

    plott(y_test, y_pred, plot_title)

    return mean_squared_error(y_test, y_pred)

```

In [875]: `def sklearn_SGD(X_train, y_train, max_iter_given, learning_rate_given, eta_given):`

```

    if learning_rate_given == 'constant':
        clf = SGDRegressor(max_iter = max_iter_given, learning_rate= learning_rate_g
    elif learning_rate_given == 'invscaling':
        clf = SGDRegressor(max_iter = max_iter_given, learning_rate= learning_rate_g
    clf.fit(X_train, y_train)
    mse = mean_squared_error(y_train, clf.predict(X_train))
    return clf, mse


```

[3.2] Applying Mini-Batch SGD with fixed learning rate    itr = 100

```
In [876]: # Using implemented Mini-Batch SGD with constant learning rate
          X_samples, Y_samples, y_pred, w_latest1, b_latest1, train_error1 = stochastic_gradier
          print(train_error1)
```

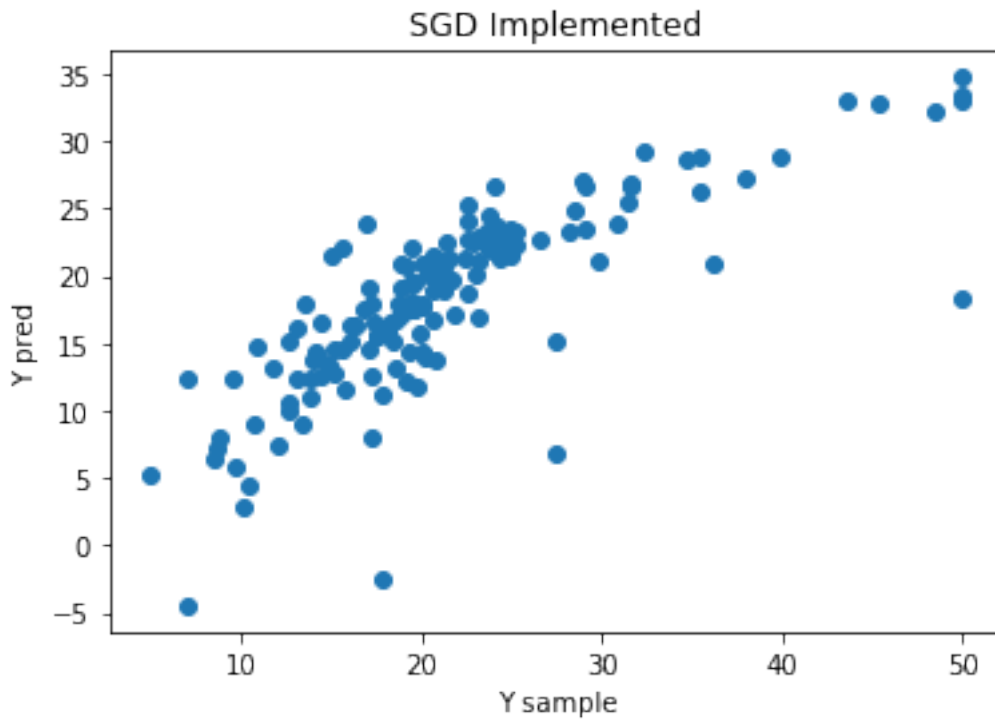
48.92878984555921

```
In [877]: plott(Y_samples, y_pred, 'SGD Implemented')
```



```
In [878]: test_error1 = mean_sq_error(X_test, y_test, w_latest1, b_latest1, 'SGD Implemented')
          print(test_error1)
```

36.732583474085736



In [879]: *# Using Sklearn*

```
clf1, train_mse1 = sklearn_SGD(X_train, y_train, 100, 'constant', 0.01)
test_mse1 = mean_squared_error(y_test, clf1.predict(X_test))
train_mse1
```


Out[879]: 25.44650397100788

itr = 1000

```
In [880]: X_samples, Y_samples, y_pred, w_latest2, b_latest2, train_error2 = stochastic_gradient_descent(X_train, y_train, 1000)
          print(train_error2)
```


26.606888825077267

```
In [881]: plott(Y_samples, y_pred, 'SGD Implemented')
```



```
In [882]: test_error2 = mean_sq_error(X_test, y_test, w_latest2, b_latest2, 'SGD Implemented')  
          print(test_error2)
```

22.161182808502716



```
In [883]: clf2, train_mse2 = sklearn_SGD(X_train, y_train, 1000, 'constant', 0.01)
          test_mse2 = mean_squared_error(y_test, clf2.predict(X_test))
          train_mse2
```


```
Out[883]: 24.546898911592436
```

```
itr = 5000
```

```
In [884]: X_samples, Y_samples, y_pred, w_latest3, b_latest3, train_error3 = stochastic_gradient_descent(X_train, y_train, 5000)
          print(train_error3)
```

```
22.69599928570044
```


```
In [885]: plott(Y_samples, y_pred, 'SGD Implemented')
```



```
In [886]: test_error3 = mean_sq_error(X_test, y_test, w_latest3, b_latest3, 'SGD Implemented')  
          print(test_error3)
```

21.264608591302853





```
In [887]: clf3, train_mse3= sklearn_SGD(X_train, y_train, 5000, 'constant', 0.01)
          test_mse3 = mean_squared_error(y_test,clf1.predict(X_test))
          train_mse3
```


```
Out[887]: 22.920499619484392
```

```
itr = 10000
```

```
In [888]: X_samples, Y_samples, y_pred, w_latest4, b_latest4, train_error4 = stochastic_gradier
          print(train_error4)
```


```
25.75845902947859
```

```
In [889]: plott(Y_samples, y_pred, 'SGD Implemented')
```



```
In [890]: test_error4 = mean_sq_error(X_test, y_test, w_latest4, b_latest4, 'SGD Implemented')  
          print(test_error4)
```

21.324201420092944



```
In [891]: clf4, train_mse4 = sklearn_SGD(X_train, y_train, 10000, 'constant', 0.01)
          test_mse4 = mean_squared_error(y_test, clf1.predict(X_test))
          train_mse4
```


```
Out[891]: 28.179224200228315
```

### [3.3] Applying Mini-Batch SGD With Constant Inverse Scaling Learning Rate itr = 100


```
In [892]: X_samples, Y_samples, y_pred, w_latest5, b_latest5, train_error5 = stochastic_gradien
          train_error5
```

```
Out[892]: 206.6501596572818
```

```
In [893]: plott(Y_samples, y_pred, 'SGD Implemented')
```



```
In [894]: test_error5 = mean_sq_error(X_test, y_test, w_latest5, b_latest5, 'SGD Implemented')
```



```
In [895]: # Using Sklearn
```

```
clf1, train_mse5 = sklearn_SGD(X_train, y_train, 100, 'invscaling', 0.01)
test_mse5 = mean_squared_error(y_test, clf1.predict(X_test))
train_mse5
```


```
Out[895]: 22.60830921541977
```

```
itr = 1000
```


```
In [896]: X_samples, Y_samples, y_pred, w_latest6, b_latest6, train_error6 = stochastic_gradient_descent(X_train, y_train, 1000)
train_error6
```

```
Out[896]: 48.2860956123789
```

```
In [897]: plott(Y_samples, y_pred, 'SGD Implemented')
```



```
In [898]: test_error6 = mean_sq_error(X_test, y_test, w_latest6, b_latest6, 'SGD Implemented')
```



In [899]: *# Using Sklearn*

```
clf1, train_mse6 = sklearn_SGD(X_train, y_train, 1000, 'invscaling', 0.01)
test_mse6 = mean_squared_error(y_test, clf1.predict(X_test))
train_mse6
```

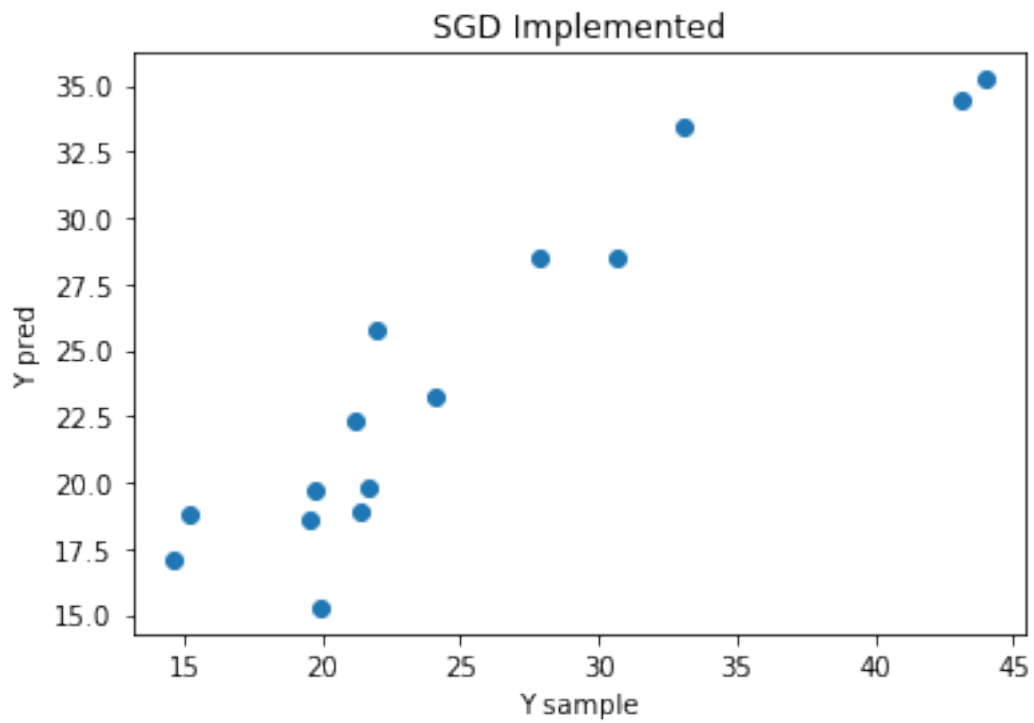
Out[899]: 22.55418955713128

itr = 5000


In [900]: X\_samples, Y\_samples, y\_pred, w\_latest7, b\_latest7, train\_error7 = stochastic\_gradie  
train\_error7

Out[900]: 14.961183281522397

In [901]: plott(Y\_samples, y\_pred, 'SGD Implemented')



```
In [902]: test_error7 = mean_sq_error(X_test, y_test, w_latest7, b_latest7, 'SGD Implemented')
```



```
In [903]: # Using Sklearn
```

```
clf1, train_mse7 = sklearn_SGD(X_train, y_train, 5000, 'invscaling', 0.01)
test_mse7 = mean_squared_error(y_test, clf1.predict(X_test))
train_mse7
```


```
Out[903]: 22.546122092391037
```

```
itr = 10000
```

```
In [904]: X_samples, Y_samples, y_pred, w_latest8, b_latest8, train_error8 = stochastic_gradie
train_error8
```


```
Out[904]: 8.308924995459867
```

```
In [905]: plott(Y_samples, y_pred, 'SGD Implemented')
```



```
In [906]: test_error8 = mean_sq_error(X_test, y_test, w_latest8, b_latest8, 'SGD Implemented')
```





```
In [907]: # Using Sklearn
          clf1, train_mse8 = sklearn_SGD(X_train, y_train, 10000, 'invscaling', 0.01)
          test_mse8 = mean_squared_error(y_test, clf1.predict(X_test))
          train_mse8
```

```
Out[907]: 22.545657510733108
```

## 1 [4.1] Conclusions

```
In [910]: #Source: http://zetcode.com/python/prettytable/
          from prettytable import PrettyTable
```

```
model_metric = PrettyTable()
```

```
model_metric = PrettyTable(["Algorithm", "Learning Rate Type", "Learning Rate", "No.
```

```
model_metric.add_row(["SGD Implemented", "constant", "0.01", "100", train_error1, tes
model_metric.add_row(["SGD Implemented", "constant", "0.01", "1000", train_error2, tes
model_metric.add_row(["SGD Implemented", "constant", "0.01", "5000", train_error3, tes
model_metric.add_row(["SGD Implemented", "constant", "0.01", "10000", train_error4, t
model_metric.add_row(["SGD Implemented", "invscaling", "0.01", "100", train_error5, t
model_metric.add_row(["SGD Implemented", "invscaling", "0.01", "1000", train_error6, t
model_metric.add_row(["SGD Implemented", "invscaling", "0.01", "5000", train_error7, t
```

```

model_metric.add_row(["SGD Implemented","invscaling", "0.01", "10000", train_error8,

model_metric.add_row(["Sklearn","constant", "0.01", "100", train_mse1, test_mse1])
model_metric.add_row(["Sklearn","constant", "0.01", "1000", train_mse2, test_mse2])
model_metric.add_row(["Sklearn","constant", "0.01", "5000", train_mse3, test_mse3])
model_metric.add_row(["Sklearn","constant", "0.01", "10000", train_mse4, test_mse4])
model_metric.add_row(["Sklearn","invscaling", "0.01", "100", train_mse5, test_mse5])
model_metric.add_row(["Sklearn","invscaling", "0.01", "1000", train_mse6, test_mse6])
model_metric.add_row(["Sklearn","invscaling", "0.01", "5000", train_mse7, test_mse7])
model_metric.add_row(["Sklearn","invscaling", "0.01", "10000", train_mse8, test_mse8])

print(model_metric.get_string(start=0))

```

Algorithm	Learning Rate Type	Learning Rate	No. of Iterations	Train MSE
SGD Implemented	constant	0.01	100	48.92878984555921
SGD Implemented	constant	0.01	1000	26.60688882507726
SGD Implemented	constant	0.01	5000	22.69599928570044
SGD Implemented	constant	0.01	10000	25.75845902947859
SGD Implemented	invscaling	0.01	100	206.6501596572818
SGD Implemented	invscaling	0.01	1000	48.2860956123789
SGD Implemented	invscaling	0.01	5000	14.96118328152239
SGD Implemented	invscaling	0.01	10000	8.308924995459867
Sklearn	constant	0.01	100	25.44650397100788
Sklearn	constant	0.01	1000	24.54689891159243
Sklearn	constant	0.01	5000	22.92049961948439
Sklearn	constant	0.01	10000	28.17922420022831
Sklearn	invscaling	0.01	100	22.60830921541977
Sklearn	invscaling	0.01	1000	22.55418955713128
Sklearn	invscaling	0.01	5000	22.54612209239103
Sklearn	invscaling	0.01	10000	22.54565751073310

## 1.1 [4.2] Steps for implementing the Mini-Batch SGD

Equations of Mini-batch SGD for Linear regression:

$$\frac{\partial L}{\partial w} = w_0 - \sum_{k=1}^j -(2/j)x_i(y_i - (w^T * x_i + b_i))$$

$$\frac{\partial L}{\partial b} = b_0 - \sum_{k=1}^j -(2/j)(y_i - (w^T * x_i + b_i))$$

- 1) Using the train data to find the components of the SGD for Linear regression such as `m_derivative(w_i)` and `b_derivative(b_i)` for a batch of randomly sampled points from train data for 'n' number of epochs. For every epoch a new batch is sampled here.

- 2) After running for n epochs, we'll obtain the optimal  $m\_derivative$  and  $b\_derivative$  values for the nth batch which is sum of  $m\_derivative$  and  $b\_derivative$  for all the points. Now multiply both of them with the learning rate to get  $w\_i$  and  $b\_i$ . Now, get  $w\_0$  and  $b\_0$  from n-1 epoch of the batch and subtract with the  $w\_i$  and  $b\_i$  respectively to get the updated values  $w\_updated$  and  $b\_updated$ .
- 3) When a point( $x\_test$ ) in test data is given we'll calculate the following for the price prediction:

$$y\_pred = w\_updated * x\_test + b\_updated$$

- 4) The above step is repeated for all the test data points and all the predicted values are stored in a list.
- 5) If the prices of the corresponding test data is present, we can plot a scatter plot for the predicted values and the actual values for the test data. Similar plot can be done for train data as well.
- 6) Mean Square Error is calculated for the test data to check the performance of the model on the unseen data. If this is high then the model performance is bad.

## 1.2 [4.3] Observations

1. Mini-Batch SGD for linear regression is implemented with the help of Sklearn's SGDRegressor. This algorithm uses two types of learning rates: a) Constant b) Inverse Scaling or inverse scaling
2. Sklearn's SGDRegressor is also utilized to generate results on the same data.
3. No. of iterations used for both type of implementations are 100, 1000, 5000, 1000
4. The Mean Square Errors(MSE) of the implemented Mini-Batch SGD both on Train and Test data are very close to that of Sklearn's SGDRegressor results except in 3 cases.
5. The results are consistent and reproducible.
6. By observation on all the graphs, the points on the graph are aligned towards a straight line.