

07 Amazon Fine Food Reviews Analysis_Support Vector Machines

March 18, 2019

1 Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective: Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative? [Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

2 [1]. Reading Data

2.1 [1.1] Loading the data

The dataset is available in two forms 1. .csv file 2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [1]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os
import sys

from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score
from sklearn.model_selection import GridSearchCV
from sklearn.calibration import CalibratedClassifierCV
from sklearn.linear_model import SGDClassifier
from sklearn.svm import SVC

In [2]: # using SQLite Table to read data.
```

```

con = sqlite3.connect(os.path.join( os.getcwd(), '..', 'database.sqlite' ))

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data points
# you can change the number to any other number based on your computing power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000 """, con)
# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 """, con)

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negative rating(-1)
def partition(x):
    if x < 3:
        return 0
    else:
        return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)

```

Number of data points in our data (525814, 10)

```

Out[2]:
   Id  ProductId  UserId  ProfileName \
0   1  B001E4KFG0  A3SGXH7AUHU8GW  delmartian
1   2  B00813GRG4  A1D87F6ZCVE5NK          dll pa
2   3  B000LQOCHO  ABXLMWJIXXAIN  Natalia Corres "Natalia Corres"

   HelpfulnessNumerator  HelpfulnessDenominator  Score  Time \
0                      1                      1      1  1303862400
1                      0                      0      0  1346976000
2                      1                      1      1  1219017600

   Summary  Text
0  Good Quality Dog Food  I have bought several of the Vitality canned d...
1    Not as Advertised  Product arrived labeled as Jumbo Salted Peanut...
2  "Delight" says it all  This is a confection that has been around a fe...

```

```

In [3]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId

```

```
HAVING COUNT(*)>1
""", con)
```

```
In [4]: print(display.shape)
display.head()
```

```
(80668, 7)
```

```
Out [4]:
```

	UserId	ProductId	ProfileName	Time	Score	\
0	#oc-R115TNMSPFT9I7	B007Y59HVM	Breyton	1331510400	2	
1	#oc-R11D9D7SHXIJB9	B005HG9ET0	Louis E. Emory "hoppy"	1342396800	5	
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1	
3	#oc-R1105J5ZVQE25C	B005HG9ET0	Penguin Chick	1346889600	5	
4	#oc-R12KPBODL2B5ZD	B0070SBE1U	Christopher P. Presta	1348617600	1	

	Text	COUNT(*)
0	Overall its just OK when considering the price...	2
1	My wife has recurring extreme muscle spasms, u...	3
2	This coffee is horrible and unfortunately not ...	2
3	This will be the bottle that you grab from the...	3
4	I didnt like this coffee. Instead of telling y...	2

```
In [5]: display[display['UserId']=='AZY10LLTJ71NX']
```

```
Out [5]:
```

	UserId	ProductId	ProfileName	Time	\
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200	

	Score	Text	COUNT(*)
80638	5	I was recommended to try green tea extract to ...	5

```
In [6]: display['COUNT(*)'].sum()
```

```
Out [6]: 393063
```

3 [2] Exploratory Data Analysis

3.1 [2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [7]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

```

Out [7]:
      Id  ProductId      UserId      ProfileName  HelpfulnessNumerator  \
0   78445  B000HDL1RQ  AR5J8UI46CURR  Geetha Krishnan                2
1  138317  B000HDOPYC  AR5J8UI46CURR  Geetha Krishnan                2
2  138277  B000HDOPYM  AR5J8UI46CURR  Geetha Krishnan                2
3   73791  B000HDOPZG  AR5J8UI46CURR  Geetha Krishnan                2
4  155049  B000PAQ75C  AR5J8UI46CURR  Geetha Krishnan                2

      HelpfulnessDenominator  Score      Time  \
0                        2      5  1199577600
1                        2      5  1199577600
2                        2      5  1199577600
3                        2      5  1199577600
4                        2      5  1199577600

                        Summary  \
0  LOACKER QUADRATINI VANILLA WAFERS
1  LOACKER QUADRATINI VANILLA WAFERS
2  LOACKER QUADRATINI VANILLA WAFERS
3  LOACKER QUADRATINI VANILLA WAFERS
4  LOACKER QUADRATINI VANILLA WAFERS

                        Text
0  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
1  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
2  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
3  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
4  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...

```

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8) ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```

In [8]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False)

In [9]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first')
final.shape

```

```
Out[9]: (364173, 10)
```

```
In [10]: #Checking to see how much % of data still remains
         (final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

```
Out[10]: 69.25890143662969
```

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```
In [11]: display= pd.read_sql_query("""
        SELECT *
        FROM Reviews
        WHERE Score != 3 AND Id=44737 OR Id=64422
        ORDER BY ProductID
        """, con)
```

```
display.head()
```

```
Out[11]:
```

	Id	ProductId	UserId	ProfileName	\
0	64422	B000MIDR0Q	A161DK06JJMCYF	J. E. Stephens	"Jeanne"
1	44737	B001EQ55RW	A2V0I904FH7ABY		Ram

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0		3	1	5	1224892800
1		3	2	4	1212883200

	Summary	\
0	Bought This for My Son at College	
1	Pure cocoa taste with crunchy almonds inside	

	Text
0	My son loves spaghetti so I didn't hesitate or...
1	It was almost a 'love at first bite' - the per...

```
In [12]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [13]: #Before starting the next phase of preprocessing lets see the number of entries left
         print(final.shape)
```

```
#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()
```

```
(364171, 10)
```

```
Out[13]: 1    307061
         0    57110
         Name: Score, dtype: int64
```

4 [3] Preprocessing

4.1 [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [14]: # printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)
```

```
this witty little book makes my son laugh at loud. i recite it in the car as we're driving along
=====
I was really looking forward to these pods based on the reviews. Starbucks is good, but I prefer
=====
Great ingredients although, chicken should have been 1st rather than chicken broth, the only thing
=====
Can't do sugar. Have tried scores of SF Syrups. NONE of them can touch the excellence of this
=====
```

```
In [15]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
```

```

sent_150 = re.sub(r"http\S+", "", sent_1500)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)

```

this witty little book makes my son laugh at loud. i recite it in the car as we're driving along

```

In [16]: # https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all
from bs4 import BeautifulSoup

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)

```

this witty little book makes my son laugh at loud. i recite it in the car as we're driving along

```

=====
I was really looking forward to these pods based on the reviews. Starbucks is good, but I prefer
=====
Great ingredients although, chicken should have been 1st rather than chicken broth, the only thing
=====
Can't do sugar. Have tried scores of SF Syrups. NONE of them can touch the excellence of this

```

```

In [17]: # https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)

```



```

phrase = re.sub(r"\ 're", " are", phrase)
phrase = re.sub(r"\ 's", " is", phrase)
phrase = re.sub(r"\ 'd", " would", phrase)
phrase = re.sub(r"\ 'll", " will", phrase)
phrase = re.sub(r"\ 't", " not", phrase)
phrase = re.sub(r"\ 've", " have", phrase)
phrase = re.sub(r"\ 'm", " am", phrase)
return phrase

```

```

In [18]: sent_1500 = decontracted(sent_1500)
print(sent_1500)
print("="*50)

```

Great ingredients although, chicken should have been 1st rather than chicken broth, the only thing
=====

```

In [19]: #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent_0 = re.sub("\S*\d\S*", "", sent_0).strip()
print(sent_0)

```

this witty little book makes my son laugh at loud. i recite it in the car as we're driving along

```

In [20]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)

```

Great ingredients although chicken should have been 1st rather than chicken broth the only thing

```

In [21]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have reuvmoved in the 1st step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves',
'you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him',
'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself',
'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', 'that',
'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had',
'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as',
'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through',
'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over',
'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any',
'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too',
's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'n',
've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't",

```

```
"hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mi',
"mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't",
'won', "won't", 'wouldn', "wouldn't"])
```

5 Applying SVM

[3.2] Preprocessing Review Text and Summary

```
In [22]: preprocessed_reviews = []
preprocessed_summary = []
# Sampling the data and preprocessing
def data_sampling_preprocessing(final, no_of_samples):

    final = final.sample(n=no_of_samples)

    # Combining all the above students
    from tqdm import tqdm
    preprocessed_reviews = []
    # tqdm is for printing the status bar
    for sentence in tqdm(final['Text'].values):
        sentence = re.sub(r"http\S+", "", sentence)
        sentence = BeautifulSoup(sentence, 'lxml').get_text()
        sentence = decontracted(sentence)
        sentence = re.sub("\S*\d\S*", "", sentence).strip()
        sentence = re.sub('[^A-Za-z]+', ' ', sentence)
        # https://gist.github.com/sebleier/554280
        sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stop)
        preprocessed_reviews.append(sentence.strip())

    # Combining all the above students
    from tqdm import tqdm
    preprocessed_summary = []
    # tqdm is for printing the status bar
    for summary in tqdm(final['Summary'].values):
        summary = re.sub(r"http\S+", "", summary)
        summary = BeautifulSoup(summary, 'lxml').get_text()
        summary = decontracted(summary)
        summary = re.sub("\S*\d\S*", "", summary).strip()
        summary = re.sub('[^A-Za-z]+', ' ', summary)
        # https://gist.github.com/sebleier/554280
        summary = ' '.join(e.lower() for e in summary.split() if e.lower() not in stop)

        preprocessed_summary.append(summary.strip())

    final['CleanedText'] = preprocessed_reviews #adding a column of CleanedText which
    final['CleanedText'] = final['CleanedText'].astype('str')
```

```

final['CleanedSummary'] = preprocessed_summary #adding a column of CleanedSummary
final['CleanedSummary'] = final['CleanedSummary'].astype('str')

final['Text_Summary'] = final['CleanedSummary'] + final['CleanedText']

# # store final table into an SQLite table for future.
# conn = sqlite3.connect('final.sqlite')
# c=conn.cursor()
# conn.text_factory = str
# final.to_sql('Reviews', conn, schema=None, if_exists='replace', \
#             index=True, index_label=None, chunksize=None, dtype=None)
# conn.close()

return final['Text_Summary'], final['Score']

```

6 [4] Featurization

6.1 [4.1] BAG OF WORDS

```

In [23]: # #BoW
# count_vect = CountVectorizer() #in scikit-learn
# count_vect.fit(preprocessed_reviews)
# print("some feature names ", count_vect.get_feature_names()[:10])
# print('='*50)

# final_counts = count_vect.transform(preprocessed_reviews)
# print("the type of count vectorizer ",type(final_counts))
# print("the shape of out text BOW vectorizer ",final_counts.get_shape())
# print("the number of unique words ", final_counts.get_shape()[1])

```

6.2 [4.2] Bi-Grams and n-Grams.

```

In [24]: # #bi-gram, tri-gram and n-gram

# #removing stop words like "not" should be avoided before building n-grams
# # count_vect = CountVectorizer(ngram_range=(1,2))
# # please do read the CountVectorizer documentation http://scikit-learn.org/stable/m

# # you can choose these numebrs min_df=10, max_features=5000, of your choice
# count_vect = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=5000)
# final_bigram_counts = count_vect.fit_transform(preprocessed_reviews)
# print("the type of count vectorizer ",type(final_bigram_counts))
# print("the shape of out text BOW vectorizer ",final_bigram_counts.get_shape())
# print("the number of unique words including both unigrams and bigrams ", final_bigr

```

6.3 [4.3] TF-IDF

```
In [25]: # tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10)
# tf_idf_vect.fit(preprocessed_reviews)
# print("some sample features(unique words in the corpus)",tf_idf_vect.get_feature_names())
# print('='*50)

# final_tf_idf = tf_idf_vect.transform(preprocessed_reviews)
# print("the type of count vectorizer ",type(final_tf_idf))
# print("the shape of out text TFIDF vectorizer ",final_tf_idf.get_shape())
# print("the number of unique words including both unigrams and bigrams ", final_tf_idf.get_shape()[0])
```

6.4 [4.4] Word2Vec

```
In [26]: # # Train your own Word2Vec model using your own text corpus
# i=0
# list_of_sentence=[]
# for sentence in preprocessed_reviews:
#     list_of_sentence.append(sentence.split())
```

```
In [27]: # # Using Google News Word2Vectors
```

```
# # in this project we are using a pretrained model by google
# # its 3.3G file, once you load this into your memory
# # it occupies ~9Gb, so please do this step only if you have >12G of ram
# # we will provide a pickle file wich contains a dict ,
# # and it contains all our courpus words as keys and model[word] as values
# # To use this code-snippet, download "GoogleNews-vectors-negative300.bin"
# # from https://drive.google.com/file/d/0B7XkCwpI5KDYNlNUTTlSS21pQmM/edit
# # it's 1.9GB in size.

# # http://kavita-ganesan.com/gensim-word2vec-tutorial-starter-code/#.W17SRFAzZPY
# # you can comment this whole cell
# # or change these variable according to your need

# is_your_ram_gt_16g=False
# want_to_use_google_w2v = False
# want_to_train_w2v = True

# if want_to_train_w2v:
#     # min_count = 5 considers only words that occurred atleast 5 times
#     w2v_model=Word2Vec(list_of_sentence,min_count=5,size=50, workers=4)
#     print(w2v_model.wv.most_similar('great'))
#     print('='*50)
#     print(w2v_model.wv.most_similar('worst'))

# elif want_to_use_google_w2v and is_your_ram_gt_16g:
#     if os.path.isfile('GoogleNews-vectors-negative300.bin'):
```

```

#         w2v_model=KeyedVectors.load_word2vec_format('GoogleNews-vectors-negative300
#         print(w2v_model.wv.most_similar('great'))
#         print(w2v_model.wv.most_similar('worst'))
#     else:
#         print("you don't have gogole's word2vec file, keep want_to_train_w2v = True

```

```

In [28]: # w2v_words = list(w2v_model.wv.vocab)
# print("number of words that occured minimum 5 times ",len(w2v_words))
# print("sample words ", w2v_words[0:50])

```

6.5 [4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V

[4.4.1.1] Avg W2v

```

In [29]: # # average Word2Vec
# # compute average word2vec for each review.
# sent_vectors = []; # the avg-w2v for each sentence/review is stored in this list
# for sent in tqdm(list_of_sentence): # for each review/sentence
#     sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need
#     cnt_words=0; # num of words with a valid vector in the sentence/review
#     for word in sent: # for each word in a review/sentence
#         if word in w2v_words:
#             vec = w2v_model.wv[word]
#             sent_vec += vec
#             cnt_words += 1
#     if cnt_words != 0:
#         sent_vec /= cnt_words
#     sent_vectors.append(sent_vec)
# print(len(sent_vectors))
# print(len(sent_vectors[0]))

```

[4.4.1.2] TFIDF weighted W2v

```

In [30]: # # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
# model = TfidfVectorizer()
# tf_idf_matrix = model.fit_transform(preprocessed_reviews)
# # we are converting a dictionary with word as a key, and the idf as a value
# dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

```

```

In [31]: # # TF-IDF weighted Word2Vec
# tfidf_feat = model.get_feature_names() # tfidf words/col-names
# # final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfi

# tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this
# row=0;
# for sent in tqdm(list_of_sentence): # for each review/sentence
#     sent_vec = np.zeros(50) # as word vectors are of zero length
#     weight_sum =0; # num of words with a valid vector in the sentence/review
#     for word in sent: # for each word in a review/sentence

```

```

#         if word in w2v_words and word in tfidf_feat:
#             vec = w2v_model.wv[word]
#             tfidf = tfidf_matrix[row, tfidf_feat.index(word)]
#             # to reduce the computation we are
#             # dictionary[word] = idf value of word in whole corpus
#             # sent.count(word) = tf value of word in this review
#             tfidf = dictionary[word]*(sent.count(word)/len(sent))
#             sent_vec += (vec * tfidf)
#             weight_sum += tfidf
#         if weight_sum != 0:
#             sent_vec /= weight_sum
#         tfidf_sent_vectors.append(sent_vec)
#         row += 1

```

7 [5] Assignment 7: SVM

Apply SVM on these feature sets

SET 1:Review text, preprocessed one converted into vectors

SET 2:Review text, preprocessed one converted into vectors

SET 3:Review text, preprocessed one converted into vectors

SET 4:Review text, preprocessed one converted into vectors

Procedure

You need to work with 2 versions of SVM

Linear kernel

RBF kernel

When you are working with linear kernel, use SGDClassifier with hinge loss because it is c

When you are working with SGDClassifier with hinge loss and trying to find the AUC

score, you would have to use <a href='https://scikit-learn.org/stable/modules/generated/sk

Similarly, like kdtree of knn, when you are working with RBF kernel it's better to reduce

the number of dimensions. You can put min_df = 10, max_features = 500 and consider a sample size of 40k points.

Hyper paramter tuning (find best alpha in range [10⁻⁴ to 10⁴], and the best pena

Find the best hyper parameter which will give the maximum <a href='https://www.appliedaicom

Find the best hyper paramter using k-fold cross validation or simple cross validation data

Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this t


```

<br>
<li><strong>Feature importance</strong>
  <ul>
<li>When you are working on the linear kernel with BOW or TFIDF please print the top 10 best
    features for each of the positive and negative classes.
  </ul>
</li>
<br>
<li><strong>Feature engineering</strong>
  <ul>
<li>To increase the performance of your model, you can also experiment with with feature engineering
    <ul>
      <li>Taking length of reviews as another feature.</li>
      <li>Considering some features from review summary as well.</li>
    </ul>
    </ul>
</li>
<br>
<li><strong>Representation of results</strong>
  <ul>
<li>You need to plot the performance of model both on train data and cross validation data for
    <img src='train_cv_auc.JPG' width=300px></li>
<li>Once after you found the best hyper parameter, you need to train your model with it, and find
    <img src='train_test_auc.JPG' width=300px></li>
<li>Along with plotting ROC curve, you need to print the <a href='https://www.appliedaicourse.com'>
    <img src='confusion_matrix.png' width=300px></li>
  </ul>
</li>
<br>
<li><strong>Conclusion</strong>
  <ul>
<li>You need to summarize the results at the end of the notebook, summarize it in the table for
    <img src='summary.JPG' width=400px>
  </li>
</ul>

```

Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakage, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method fit_transform() on you train data, and apply the method transform() on cv/test data.
4. For more details please go through this link.

8 Applying SVM

In [32]: *# Source: <https://docs.python.org/3/library/pickle.html>*

```
# Saving data to pickle file
def topicklefile(obj, file_name):
    pickle.dump(obj,open(file_name+'.pkl', 'wb'))
```

In [33]: *# Data from pickle file*

```
def frompicklefile(file_name):
    data = pickle.load(open(file_name+'.pkl', 'rb'))
    return data
```

In [34]: *# Sort 'Time' column*

```
final = final.sort_values(by='Time', ascending=True)
```

In [35]: *# Source: https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html*

```
# Train Test split for train and test data
def data_split(final, no_of_samples):
    X, y = data_sampling_preprocessing(final, no_of_samples)
    # split the data set into train and test
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
    topicklefile(X_train, 'X_train')
    topicklefile(X_test, 'X_test')
    topicklefile(y_train, 'y_train')
    topicklefile(y_test, 'y_test')

    return X_train, X_test, y_train, y_test
```

In [36]: `def apply_avgw2v_train_test(X_train, X_test):`

```
# Training own Word2Vec model using your own text corpus
list_of_sent_train = []
for sent in X_train:#final['Text_Summary'].values:
    list_of_sent_train.append(sent.split())
list_of_sent_test = []
for sent in X_test:#final['Text_Summary'].values:
    list_of_sent_test.append(sent.split())

# min_count = 5 considers only words that occurred atleast 5 times
w2v_model=Word2Vec(list_of_sent_train,min_count=5,size=50, workers=8)

w2v_words = list(w2v_model.wv.vocab)
# print("number of words that occurred minimum 5 times ",len(w2v_words))
# print("sample words ", w2v_words[0:50])

# compute average word2vec for each review for train data
avgw2v_train = []; # the avg-w2v for each sentence/review is stored in this list
```



```

for sent in tqdm(list_of_sent_train): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    avgw2v_train.append(sent_vec)
# print(len(avgw2v_train))
# print(len(avgw2v_train[0]))

# compute average word2vec for each review for test data
avgw2v_test = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_sent_test): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    avgw2v_test.append(sent_vec)
# print(len(avgw2v_test))
# print(len(avgw2v_test[0]))

return avgw2v_train, avgw2v_test

```

In [37]: `def apply_tfidfw2v_train_test(X_train, X_test):`

```

# Training own Word2Vec model using your own text corpus
list_of_sent_train = []
for sent in X_train:#final['Text_Summary'].values:
    list_of_sent_train.append(sent.split())
list_of_sent_test = []
for sent in X_test:#final['Text_Summary'].values:
    list_of_sent_test.append(sent.split())

# min_count = 5 considers only words that occurred atleast 5 times
w2v_model=Word2Vec(list_of_sent_train,min_count=5,size=50, workers=8)

w2v_words = list(w2v_model.wv.vocab)

```

```

model = TfidfVectorizer()
tfidf_matrix = model.fit_transform(X_train)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

# TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = t

tfidf2v_train = []; # the tfidf-w2v for each sentence/review is stored in this li
row=0;
for sent in tqdm(list_of_sent_train): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
#            tf_idf = tfidf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf2v_train.append(sent_vec)
    row += 1

tfidf_matrix = model.transform(X_test)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

# TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = t

tfidf2v_test = []; # the tfidf-w2v for each sentence/review is stored in this li
row=0;
for sent in tqdm(list_of_sent_test): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
#            tf_idf = tfidf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are

```

```

        # dictionary[word] = idf value of word in whole corpus
        # sent.count(word) = tf value of word in this review
        tf_idf = dictionary[word]*(sent.count(word)/len(sent))
        sent_vec += (vec * tf_idf)
        weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf2v_test.append(sent_vec)
    row += 1

return tfidf2v_train, tfidf2v_test

```

```

In [38]: # Applying BOW on train and test data and creating the
from sklearn.preprocessing import StandardScaler
from scipy.sparse import hstack

```

```

#Standardize 'bow_train' data features by removing the mean and scaling to unit variance
std_scalar1 = StandardScaler(copy=True, with_mean=False, with_std=True)
std_scalar2 = StandardScaler(copy=True, with_mean=True, with_std=True)

```

```
count = 0
```

```
def apply_vectorizers_train_test(final, algo, model_name):
```

```

    global count
    if count == 0 or count == 4:
        if algo == 'LinearSVM':
            train_data, test_data, y_train, y_test = data_split(final, 100000)
        elif algo == 'RBFKernel':
            train_data, test_data, y_train, y_test = data_split(final, 40000)
        count += 1
        print("count: ", count)
        topicklefile(train_data, 'train_data')
        topicklefile(test_data, 'test_data')
        topicklefile(y_train, 'y_train')
        topicklefile(y_test, 'y_test')
    else:
        train_data = frompicklefile('train_data')
        test_data = frompicklefile('test_data')
        y_train = frompicklefile('y_train')
        y_test = frompicklefile('y_test')
        count += 1
        print("count: ", count)

    if model_name == 'BOW':
        #Applying BoW on Train data

```

```

if algo == 'LinearSVM':
    count_vect = CountVectorizer()
elif algo == 'RBFKernel':
    count_vect = CountVectorizer(min_df = 10, max_features = 500)

#Applying BoW on Test data
bow_train_vect = count_vect.fit_transform(train_data)

#Applying BoW on Test data similar to the bow_train data
bow_test_vect = count_vect.transform(test_data)

# Standardise train data
bow_train_vect = std_scalar1.fit_transform(bow_train_vect)
# Standardize the unseen bow_test data
bow_test_vect = std_scalar1.transform(bow_test_vect)

topicklefile(bow_train_vect, 'bow_train_vect')
topicklefile(bow_test_vect, 'bow_test_vect')

print("'bow_train_vect' and 'bow_test_vect' are the pickle files.")
return count_vect

elif model_name == 'TF-IDF':
    #Applying TF-IDF on Train data
    if algo == 'LinearSVM':
        count_vect = TfidfVectorizer(ngram_range=(1,2))
    elif algo == 'RBFKernel':
        count_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10, max_features =

    #Applying BoW on Test data
    tfidf_train_vect = count_vect.fit_transform(train_data)

    #Applying BoW on Test data similar to the bow_train data
    tfidf_test_vect = count_vect.transform(test_data)

    # Standardise train data
    tfidf_train_vect = std_scalar1.fit_transform(tfidf_train_vect)
    # Standardize the unseen bow_test data
    tfidf_test_vect = std_scalar1.transform(tfidf_test_vect)

    topicklefile(tfidf_train_vect, 'tfidf_train_vect')
    topicklefile(tfidf_test_vect, 'tfidf_test_vect')

    print("'tfidf_train_vect' and 'tfidf_test_vect' are the pickle files.")
    return count_vect

elif model_name == 'AvgW2V':
    avgw2v_train_vect, avgw2v_test_vect = apply_avgw2v_train_test(train_data, test_data)

```

```

    # Standardise train data
    avgw2v_train_vect = std_scalar2.fit_transform(avgw2v_train_vect)
    # Standardize the unseen bow_test data
    avgw2v_test_vect = std_scalar2.transform(avgw2v_test_vect)

    topicklefile(train_vect, 'avgw2v_train_vect')
    topicklefile(test_vect, 'avgw2v_test_vect')
    print("'avgw2v_train_vect' and 'avgw2v_test_vect' are the pickle files.")

elif model_name == 'TF-IDF W2V':
    tfidf2v_train_vect, tfidf2v_test_vect = apply_tfidf2v_train_test(train_data, y_train)

    # Standardise train data
    tfidf2v_train_vect = std_scalar2.fit_transform(tfidf2v_train_vect)
    # Standardize the unseen bow_test data
    tfidf2v_test_vect = std_scalar2.transform(tfidf2v_test_vect)

    topicklefile(train_vect, 'tfidf2v_train_vect')
    topicklefile(test_vect, 'tfidf2v_test_vect')
    print("'tfidf2v_train_vect' and 'tfidf2v_test_vect' are the pickle files.")

else:
    #Error Message
    print('Model specified is not valid! Please check.')

```

```
In [39]: def apply_svm(algo, parameters, train_data, y_train):
```

```

    if algo == 'LinearSVM':
        #         parameters = {'alpha':alpha_values}
        svm_clf = SGDClassifier(class_weight='balanced')

        clf = GridSearchCV(svm_clf, parameters, cv=10, scoring= 'roc_auc', n_jobs=-1,
        clf.fit(train_data, y_train)
        alpha_optimal = clf.best_params_.get('alpha')

        optimal_hyperparameter = alpha_optimal
        optimal_penalty= clf.best_params_.get('penalty')

        #Getting the Train and CV AUC score values for only 'optimal_penalty' for the
        clf_cv_results = pd.DataFrame(clf.cv_results_)
        clf_cv_results = clf_cv_results[clf_cv_results['param_penalty']== optimal_penalty]

        train_auc= clf_cv_results['mean_train_score']
        train_auc_std= clf_cv_results['std_train_score']
        cv_auc = clf_cv_results['mean_test_score']
        cv_auc_std= clf_cv_results['std_test_score']

```

```

        return clf, optimal_penalty, optimal_hyperparameter, train_auc, train_auc_std,

    elif algo == 'RBFKernel':
#         parameters = {'C':alpha_values}
        svm_clf = SVC(kernel='rbf', probability = True, class_weight='balanced', cache=1000)

        clf = GridSearchCV(svm_clf, parameters, cv=10, scoring= 'roc_auc', n_jobs=-1,
                           verbose=1)
        clf.fit(train_data, y_train)
        c_optimal = clf.best_params_.get('C')

        optimal_hyperparameter = c_optimal

#         print(clf.cv_results_)

        train_auc= clf.cv_results_['mean_train_score']
        train_auc_std= clf.cv_results_['std_train_score']
        cv_auc = clf.cv_results_['mean_test_score']
        cv_auc_std= clf.cv_results_['std_test_score']

        return clf, optimal_hyperparameter, train_auc, train_auc_std, cv_auc, cv_auc_std

In [40]: def train_cv_error_plot(alpha_values, train_auc, train_auc_std, cv_auc, cv_auc_std):
#         alpha_values = np.log10(alpha_values)
        plt.plot(np.log10(alpha_values), train_auc, label='Train AUC')

#         # Source: https://stackoverflow.com/a/48803361/4084039
#         plt.gca().fill_between(alpha_values,train_auc - train_auc_std,train_auc + train_auc_std)

        plt.plot(np.log10(alpha_values), cv_auc, label='CV AUC')
#         # Source: https://stackoverflow.com/a/48803361/4084039
#         plt.gca().fill_between(alpha_values,cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha_values)
        plt.legend()
        plt.xlabel("Hyperparameter-log Values")
        plt.ylabel("AUC")
        plt.title("ERROR PLOTS")
        plt.show()

In [41]: def svm_optimal(algo, optimal_hyperparameter, optimal_penalty, train_vect, y_train):
    if algo == 'LinearSVM':
        optimal_svm = SGDClassifier(loss="hinge",penalty = optimal_penalty, alpha = 0.0001)
        optimal_svm.fit(train_vect, y_train)
    elif algo == 'RBFKernel':
        optimal_svm = SVC(kernel='rbf', C = optimal_hyperparameter, probability = True)
        optimal_svm.fit(train_vect, y_train)
    return optimal_svm

In [42]: def retrain_svm(optimal_svm, train_vect, y_train, test_vect, y_test):

```

```

# fitting the model with optimal K for training data
optimal_svm.fit(train_vect, y_train)

# predict the response for the unseen bow_test data
y_pred = optimal_svm.predict(test_vect)

In [43]: # Confusion Matrix
def cm_fig(optimal_svm, y_test, test_vect):
    cm = pd.DataFrame(confusion_matrix(y_test, optimal_svm.predict(test_vect)))
    print(confusion_matrix(y_test, optimal_svm.predict(test_vect)))

    plt.figure(1, figsize=(18,5))
    plt.subplot(121)
    plt.title("Confusion Matrix")
    sns.set(font_scale=1.4)
    sns.heatmap(cm, cmap= 'gist_earth', annot=True, annot_kws={'size':15}, fmt='g')

In [44]: def svm_calibratedclassifierCV(algo, optimal_svm, penalty_given, train_data, y_train)
    if algo == 'LinearSVM':
        svm_calib = CalibratedClassifierCV(optimal_svm, method = 'sigmoid', cv='prefi
        svm_calib.fit(train_data, y_train)
    elif algo == 'RBFKernel':
        print("No need of Calibrated Classifier CV for SVC")
        svm_calib = optimal_svm
    #         svm_calib = CalibratedClassifierCV(optimal_svm, method='sigmoid', cv=5)
    #         svm_calib.fit(train_data, y_train)
    return svm_calib

In [45]: #Reference: https://stackoverflow.com/questions/52910061/implementing-roc-curves-for-
def error_plot(svm_obj, train_vec, y_train, test_vec, y_test):
    train_fpr, train_tpr, thresholds = roc_curve(y_train, svm_obj.predict_proba(train_vec)[:,1])
    test_fpr, test_tpr, thresholds = roc_curve(y_test, svm_obj.predict_proba(test_vec)[:,1])

    plt.plot(train_fpr, train_tpr, label="train AUC = %0.3f" %auc(train_fpr, train_tpr))
    plt.plot(test_fpr, test_tpr, label="train AUC = %0.3f" %auc(test_fpr, test_tpr))
    plt.plot([0.0, 1.0], [0.0, 1.0], 'k--')
    plt.legend()
    plt.xlabel("False Positive Rate")
    plt.ylabel("True Positive Rate")
    plt.title("ROC Curve")
    plt.show()

    return auc(test_fpr, test_tpr)

In [46]: def get_features_top(count_vect, optimal_svm):
    features=count_vect.get_feature_names()
    feature_prob=optimal_svm.coef_.ravel()
    print(len(features))
    print('='*100)

```

```

print(feature_prob.shape)
df_feature_proba = pd.DataFrame({'features':features, 'probabilities':feature_prob})
df_feature_proba = df_feature_proba.sort_values(by=['probabilities'],ascending=False)
# print(df_feature_proba)
return df_feature_proba[:11], df_feature_proba[-11:]

```

8.1 [5.1] Linear SVM

8.1.1 [5.1.1] Applying Linear SVM on BOW, SET 1

In [47]: *# Please write all the code with proper documentation*

In [48]: bow_count_vect = apply_vectorizers_train_test(final, 'LinearSVM', 'BOW')

```

100%|| 100000/100000 [00:41<00:00, 2395.96it/s]
100%|| 100000/100000 [00:26<00:00, 3724.77it/s]

```

```

count: 1
'bow_train_vect' and 'bow_test_vect' are the pickle files.

```

```

In [49]: train_vect = frompicklefile('bow_train_vect')
test_vect = frompicklefile('bow_test_vect')
y_train = frompicklefile('y_train')
y_test = frompicklefile('y_test')

```

```

print(train_vect.shape)
print(len(y_train))

```

```

(70000, 100497)
70000

```

```

In [50]: alpha_values = [10**-5, 10**-4, 10**-3, 10**-2, 10**-1, 1, 10**1, 10**2, 10**3, 10**4]
penalties = ['l1', 'l2']
# penalties = ['l1']
hyper_parameters = {'alpha':alpha_values, 'penalty':penalties}
clf, optimal_penalty, optimal_hyperparameter, train_auc, train_auc_std, cv_auc, cv_auc_std = grid_search(hyper_parameters)

print('The optimal penalty is {}'.format(optimal_penalty))
print('The optimal hyperparameter is {}'.format(optimal_hyperparameter))

```

```

optimal_hyperparameter_bow1 = optimal_hyperparameter


```

```

The optimal penalty is l2
The optimal hyperparameter is 10

```

In [51]: train_cv_error_plot(alpha_values, train_auc, train_auc_std, cv_auc, cv_auc_std)

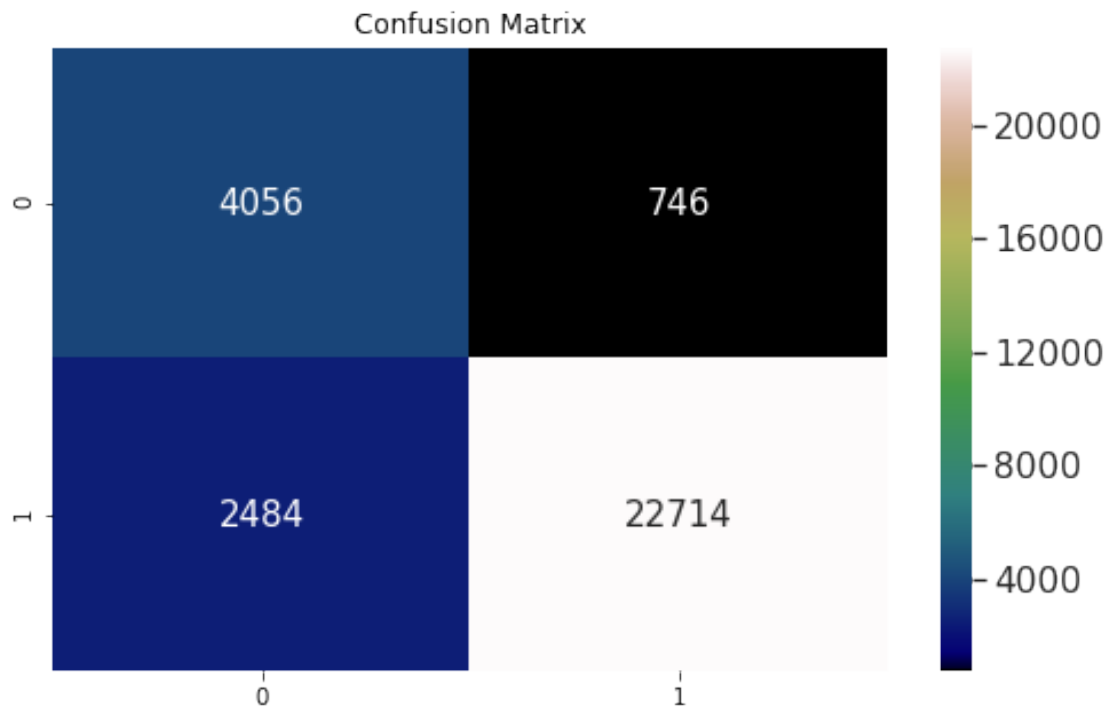


```
In [52]: bow_optimal_svm = svm_optimal('LinearSVM', optimal_hyperparameter_bow1, optimal_penal
```


```
In [53]: # retrain_sum(bow_optimal_svm, train_vect, y_train, test_vect, y_test)
```

```
In [54]: cm_fig(bow_optimal_svm, y_test, test_vect)
```

```
[[ 4056   746]
 [ 2484 22714]]
```



```
In [55]: svm_calib = svm_calibratedclassifierCV('LinearSVM', bow_optimal_svm, optimal_penalty,  
In [56]: bow_auc1 = error_plot(svm_calib, train_vect, y_train, test_vect, y_test)  
          bow_auc1
```



Out [56]: 0.9386344698096035

In [57]: positive_features, negative_features = get_features_top(bow_count_vect, bow_optimal_score, positive_features)

100497

(100497,)

Out [57]:

	features	probabilities
40073	great	0.023859
7471	best	0.015814
52053	love	0.013308
23230	delicious	0.013134
64990	perfect	0.011331
38723	good	0.010757
30778	excellent	0.010381
32646	favorite	0.009903
43189	highly	0.009863
52616	loves	0.009572
97611	wonderful	0.009149

```
In [58]: negative_features
```

```
Out[58]:
```

	features	probabilities
25521	disappointing	-0.012450
89812	threw	-0.012903
5092	awful	-0.013096
43949	horrible	-0.013649
57159	money	-0.013701
88823	terrible	-0.013801
5450	bad	-0.014145
98236	worst	-0.014452
95998	waste	-0.014561
25409	disappointed	-0.015801
59781	not	-0.019110

8.1.2 [5.1.2] Applying Linear SVM on TFIDF, SET 2

```
In [59]: # Please write all the code with proper documentation
```

```
In [60]: tfidf_count_vect = apply_vectorizers_train_test(final, 'LinearSVM', 'TF-IDF')
```

```
count: 2
```

'tfidf_train_vect' and 'tfidf_test_vect' are the pickle files.

```
In [61]: train_vect = frompicklefile('tfidf_train_vect')
test_vect = frompicklefile('tfidf_test_vect')
y_train = frompicklefile('y_train')
y_test = frompicklefile('y_test')
```

```
print(train_vect.shape)
print(len(y_train))
```

```
(70000, 1384302)
```

```
70000
```

```
In [62]: alpha_values = [10**-5, 10**-4, 10**-3, 10**-2, 10**-1, 1, 10**1, 10**2, 10**3, 10**4, 10**5]
penalties = ['l1', 'l2']
hyper_parameters = {'alpha': alpha_values, 'penalty': penalties}
clf, optimal_penalty, optimal_hyperparameter, train_auc, train_auc_std, cv_auc, cv_auc_std = grid_search(clf, hyper_parameters, train_data, test_data)
```


```
print('The optimal penalty is {}'.format(optimal_penalty))
print('The optimal hyperparameter is {}'.format(optimal_hyperparameter))
```

```
optimal_hyperparameter_tfidf1 = optimal_hyperparameter
```

```
The optimal penalty is l2
```

```
The optimal hyperparameter is 100
```

```
In [63]: train_cv_error_plot(alpha_values,train_auc, train_auc_std, cv_auc, cv_auc_std)
```



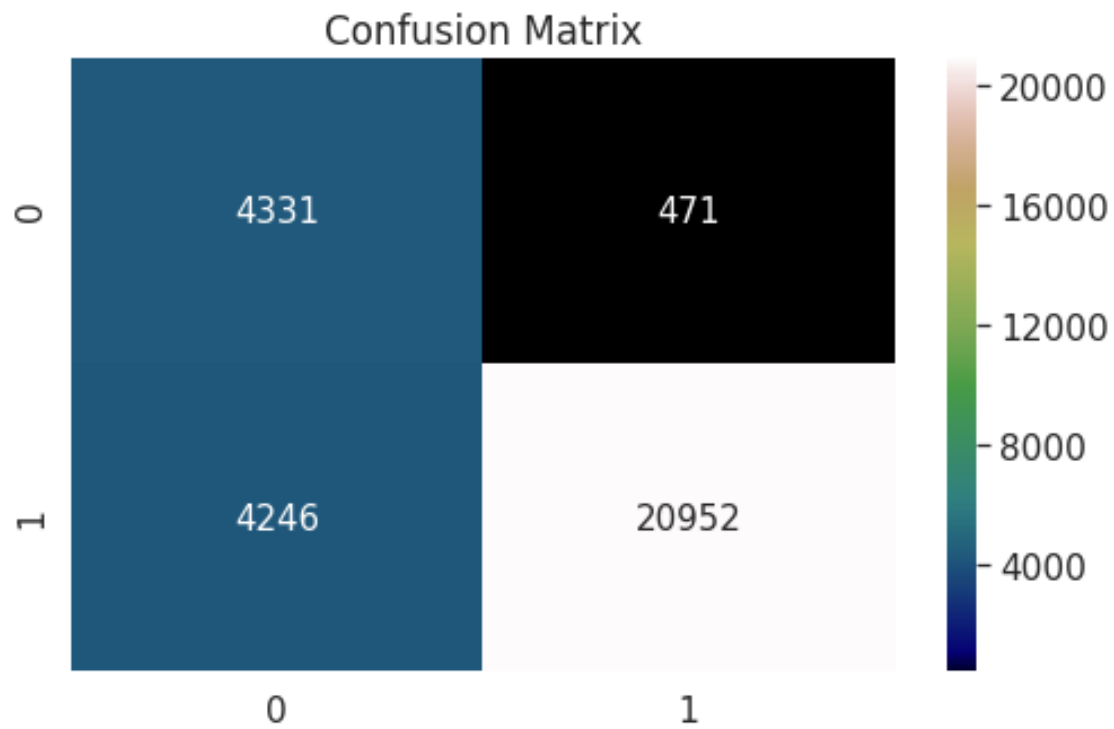
```
In [64]: tfidf_optimal_svm = svm_optimal('LinearSVM', optimal_hyperparameter_tfidf1, optimal_p
print(train_vect.shape)
print(len(y_train))
```

```
(70000, 1384302)
70000
```


```
In [65]: # retrain_svm(tfidf_optimal_svm, train_vect, y_train, test_vect, y_test)
```

```
In [66]: cm_fig(tfidf_optimal_svm, y_test, test_vect)
```

```
[[ 4331   471]
 [ 4246 20952]]
```



```
In [67]: svm_calib = svm_calibratedclassifierCV('LinearSVM', tfidf_optimal_svm, optimal_penalt  
In [68]: tfidf_auc1 = error_plot(svm_calib, train_vect, y_train, test_vect, y_test)  
          tfidf_auc1
```



Out [68]: 0.9448584206008033

In [69]: positive_features, negative_features = get_features_top(tfidf_count_vect, tfidf_optim
positive_features

1384302

(1384302,)

Out [69]:

	features	probabilities
538649	great	0.002450
100832	best	0.001634
706396	love	0.001458
313573	delicious	0.001358
521922	good	0.001228
893684	perfect	0.001173
431808	favorite	0.001054
406735	excellent	0.001017
713427	loves	0.001005
581282	highly	0.000997
367772	easy	0.000983

```
In [70]: negative_features
```

```
Out[70]:
```

	features	probabilities
1242115	threw	-0.001668
590461	horrible	-0.001733
1225716	terrible	-0.001737
773720	money	-0.001743
336662	disappointed	-0.001748
72909	bad	-0.001844
1329463	waste money	-0.001876
1364483	worst	-0.001899
1329302	waste	-0.001947
813738	not buy	-0.002005
812910	not	-0.003051

8.1.3 [5.1.3] Applying Linear SVM on AVG W2V, SET 3

```
In [71]: # Please write all the code with proper documentation
```

```
In [72]: count_vect = apply_vectorizers_train_test(final, 'LinearSVM', 'AvgW2V')
```

```
count: 3
```

```
100%|| 70000/70000 [20:10<00:00, 57.81it/s]
```

```
100%|| 30000/30000 [08:19<00:00, 60.09it/s]
```

'avgw2v_train_vect' and 'avgw2v_test_vect' are the pickle files.

```
In [73]: train_vect = frompicklefile('avgw2v_train_vect')
test_vect = frompicklefile('avgw2v_test_vect')
y_train = frompicklefile('y_train')
y_test = frompicklefile('y_test')
```

```
In [74]: alpha_values = [10**-5, 10**-4, 10**-3, 10**-2, 10**-1, 1, 10**1, 10**2, 10**3, 10**4,
penalties = ['l1', 'l2']
hyper_parameters = {'alpha': alpha_values, 'penalty': penalties}
clf, optimal_penalty, optimal_hyperparameter, train_auc, train_auc_std, cv_auc, cv_auc_std = grid_search(hyper_parameters, train_vect, test_vect, y_train, y_test)

print('The optimal penalty is {}'.format(optimal_penalty))
print('The optimal hyperparameter is {}'.format(optimal_hyperparameter))


optimal_hyperparameter_avgw2v1 = optimal_hyperparameter
```

```
The optimal penalty is l2
```

```
The optimal hyperparameter is 100
```



```
In [75]: train_cv_error_plot(alpha_values,train_auc, train_auc_std, cv_auc, cv_auc_std)
```




```
In [76]: optimal_svm = svm_optimal('LinearSVM', optimal_hyperparameter, optimal_penalty, train
```

```
In [77]: # retrain_svm(optimal_svm, train_vect, y_train, test_vect, y_test)
```


```
In [78]: cm_fig(optimal_svm, y_test, test_vect)
```

```
[[ 4337   465]
 [ 4257 20941]]
```



```
In [79]: svm_calib = svm_calibratedclassifierCV('LinearSVM', optimal_svm, optimal_penalty, tra
```

```
In [80]: avgw2v_auc1 = error_plot(svm_calib, train_vect, y_train, test_vect, y_test)
avgw2v_auc1
```



Out [80]: 0.9450313781406859

8.1.4 [5.1.4] Applying Linear SVM on TFIDF W2V, SET 4

In [81]: *# Please write all the code with proper documentation*

In [82]: `count_vect = apply_vectorizers_train_test(final, 'LinearSVM', 'TF-IDF W2V')`

count: 4

100%| 70000/70000 [1:39:08<00:00, 11.77it/s]

100%| 30000/30000 [46:50<00:00, 10.68it/s]

'tfidf_w2v_train_vect' and 'tfidf_w2v_test_vect' are the pickle files.

```
In [83]: train_vect = frompicklefile('tfidf_train_vect')
         test_vect = frompicklefile('tfidf_test_vect')
         y_train = frompicklefile('y_train')
         y_test = frompicklefile('y_test')
```


```
In [84]: alpha_values = [10**-5, 10**-4, 10**-3, 10**-2, 10**-1, 1, 10**1, 10**2, 10**3, 10**4,
penalties = ['l1', 'l2']
hyper_parameters = {'alpha':alpha_values, 'penalty':penalties}
clf, optimal_penalty, optimal_hyperparameter, train_auc, train_auc_std, cv_auc, cv_auc_std

print('The optimal penalty is {}'.format(optimal_penalty))
print('The optimal hyperparameter is {}'.format(optimal_hyperparameter))

optimal_hyperparameter_tfidf2v1 = optimal_hyperparameter
```


The optimal penalty is l2
The optimal hyperparameter is 100

```
In [85]: train_cv_error_plot(alpha_values, train_auc, train_auc_std, cv_auc, cv_auc_std)
```




```
In [86]: optimal_svm = svm_optimal('LinearSVM', optimal_hyperparameter, optimal_penalty, train_auc,
In [87]: # retrain_svm(optimal_svm, train_vect, y_train, test_vect, y_test)
In [88]: cm_fig(optimal_svm, y_test, test_vect)

[[ 4433   369]
 [ 5063 20135]]
```



```
In [89]: svm_calib = svm_calibratedclassifierCV('LinearSVM', optimal_svm, optimal_penalty, tra
```

```
In [90]: tfidf2v_auc1 = error_plot(svm_calib, train_vect, y_train, test_vect, y_test)
         tfidf2v_auc1
```



Out [90]: 0.9450462375470654

8.2 [5.2] RBF SVM

8.2.1 [5.2.1] Applying RBF SVM on BOW, SET 1

In [91]: *# Please write all the code with proper documentation*

In [92]: `count_vect = apply_vectorizers_train_test(final, 'RBFKernel', 'BOW')`

100%| 40000/40000 [00:16<00:00, 2406.48it/s]

100%| 40000/40000 [00:10<00:00, 3707.13it/s]

count: 5

'bow_train_vect' and 'bow_test_vect' are the pickle files.

In [93]: `train_vect = frompicklefile('bow_train_vect')
test_vect = frompicklefile('bow_test_vect')
y_train = frompicklefile('y_train')
y_test = frompicklefile('y_test')`


```
In [94]: C_values = [10**-5,10**-4, 10**-3, 10**-2, 10**-1, 1, 10**1, 10**2, 10**3, 10**4,10**5]
hyper_parameters = {'C':C_values}
clf, optimal_hyperparameter, train_auc, train_auc_std, cv_auc, cv_auc_std = apply_svm

print('The optimal hyperparameter is {}'.format(optimal_hyperparameter))

optimal_hyperparameter_bow2 = optimal_hyperparameter
```

The optimal hyperparameter is 0.1

```
In [95]: train_cv_error_plot(C_values,train_auc, train_auc_std, cv_auc, cv_auc_std)
```




```
In [96]: optimal_svm = svm_optimal('RBFKernel', optimal_hyperparameter, 'l2', train_vect, y_train)
```


```
In [97]: # retrain_svm(optimal_svm, train_vect, y_train, test_vect, y_test)
```

```
In [98]: cm_fig(optimal_svm, y_test, test_vect)
```

```
[[1630 224]
 [2303 7843]]
```



```
In [99]: bow_auc2 = error_plot(optimal_svm, train_vect, y_train, test_vect, y_test)
         bow_auc2
```

Out [99]: 0.9131805892863865

8.2.2 [5.2.2] Applying RBF SVM on TFIDF, SET 2

In [100]: *# Please write all the code with proper documentation*

In [101]: `count_vect = apply_vectorizers_train_test(final, 'RBFKernel', 'TF-IDF')`

count: 6

'tfidf_train_vect' and 'tfidf_test_vect' are the pickle files.

```
In [102]: train_vect = frompicklefile('tfidf_train_vect')
          test_vect = frompicklefile('tfidf_test_vect')
          y_train = frompicklefile('y_train')
          y_test = frompicklefile('y_test')
```


```
In [103]: C_values = [10**-5, 10**-4, 10**-3, 10**-2, 10**-1, 1, 10**1, 10**2, 10**3, 10**4, 10**5]
          hyper_parameters = {'C': C_values}
          clf, optimal_hyperparameter, train_auc, train_auc_std, cv_auc, cv_auc_std = apply_svm
```

```
print('The optimal hyperparameter is {}'.format(optimal_hyperparameter))
```

```
optimal_hyperparameter_tfidf2 = optimal_hyperparameter
```

The optimal hyperparameter is 1

```
In [104]: train_cv_error_plot(C_values, train_auc, train_auc_std, cv_auc, cv_auc_std)
```




```
In [105]: optimal_svm = svm_optimal('RBFKernel', optimal_hyperparameter, 'l2', train_vect, y_train)
```


```
In [106]: # retrain_svm(optimal_svm, train_vect, y_train, test_vect, y_test)
```

```
In [107]: cm_fig(optimal_svm, y_test, test_vect)
```

```
[[1331  523]
 [ 841 9305]]
```



```
In [108]: tfidf_auc2 = error_plot(optimal_svm, train_vect, y_train, test_vect, y_test)
          tfidf_auc2
```



Out[108]: 0.9235218931964408

8.2.3 [5.2.3] Applying RBF SVM on AVG W2V, SET 3

In [109]: *# Please write all the code with proper documentation*

In [110]: count_vect = apply_vectorizers_train_test(final, 'RBFKernel', 'AvgW2V')

count: 7

100%| 28000/28000 [04:18<00:00, 108.38it/s]

100%| 12000/12000 [01:53<00:00, 105.48it/s]

'avgw2v_train_vect' and 'avgw2v_test_vect' are the pickle files.

```
In [111]: train_vect = frompicklefile('avgw2v_train_vect')
          test_vect = frompicklefile('avgw2v_test_vect')
          y_train = frompicklefile('y_train')
          y_test = frompicklefile('y_test')
```


```
In [112]: C_values = [10**-5,10**-4, 10**-3, 10**-2, 10**-1, 1, 10**1, 10**2, 10**3, 10**4,10**5]
hyper_parameters = {'C':C_values}
clf, optimal_hyperparameter, train_auc, train_auc_std, cv_auc, cv_auc_std = apply_svm

print('The optimal hyperparameter is {}'.format(optimal_hyperparameter))

optimal_hyperparameter_avgw2v2 = optimal_hyperparameter
```

The optimal hyperparameter is 1

```
In [113]: train_cv_error_plot(C_values,train_auc, train_auc_std, cv_auc, cv_auc_std)
```




```
In [114]: optimal_svm = svm_optimal('RBFKernel', optimal_hyperparameter, '12', train_vect, y_train)
```


```
In [115]: # retrain_svm(optimal_svm, train_vect, y_train, test_vect, y_test)
```

```
In [116]: cm_fig(optimal_svm, y_test, test_vect)
```

```
[[1331  523]
 [ 841 9305]]
```



```
In [117]: avgw2v_auc2 = error_plot(optimal_svm, train_vect, y_train, test_vect, y_test)
          avgw2v_auc2
```



Out[117]: 0.9235226374543317

8.2.4 [5.2.4] Applying RBF SVM on TFIDF W2V, SET 4

In [118]: *# Please write all the code with proper documentation*

In [119]: `count_vect = apply_vectorizers_train_test(final, 'RBFKernel', 'TF-IDF W2V')`

count: 8

100%| 28000/28000 [17:43<00:00, 26.33it/s]

100%| 12000/12000 [07:42<00:00, 25.93it/s]

'tfidf2v_train_vect' and 'tfidf2v_test_vect' are the pickle files.

```
In [120]: train_vect = frompicklefile('tfidf2v_train_vect')
          test_vect = frompicklefile('tfidf2v_test_vect')
          y_train = frompicklefile('y_train')
          y_test = frompicklefile('y_test')
```


```
In [121]: C_values = [10**-5,10**-4, 10**-3, 10**-2, 10**-1, 1, 10**1, 10**2, 10**3, 10**4,10**5]
hyper_parameters = {'C':C_values}
clf, optimal_hyperparameter, train_auc, train_auc_std, cv_auc, cv_auc_std = apply_svm

print('The optimal hyperparameter is {}'.format(optimal_hyperparameter))

optimal_hyperparameter_tfidf2v2 = optimal_hyperparameter
```

The optimal hyperparameter is 1

```
In [122]: train_cv_error_plot(C_values,train_auc, train_auc_std, cv_auc, cv_auc_std)
```




```
In [123]: optimal_svm = svm_optimal('RBFKernel', optimal_hyperparameter, '12', train_vect, y_train)
```


```
In [124]: # retrain_svm(optimal_svm, train_vect, y_train, test_vect, y_test)
```

```
In [125]: cm_fig(optimal_svm, y_test, test_vect)
```

```
[[1331  523]
 [ 841 9305]]
```

```
In [126]: tfidf2v_auc2 = error_plot(optimal_svm, train_vect, y_train, test_vect, y_test)
          tfidf2v_auc2
```



Out[126]: 0.9235230627445552

9 [6] Conclusions

In [127]: *# Please compare all your models using Prettytable library*

In [128]: `from prettytable import PrettyTable`

```
model_metric = PrettyTable()
```

```
model_metric = PrettyTable(["Model Name", "SVM Type", 'Hyperparameter', 'AUC'])
```

```
model_metric.add_row(["Bag of Words","Linear kernel", optimal_hyperparameter_bow1, bow_1_auc])
model_metric.add_row(["TF-IDF","Linear kernel", optimal_hyperparameter_tfidf1, tfidf_1_auc])
model_metric.add_row(["Avg W2V","Linear kernel", optimal_hyperparameter_avgw2v1, avgw2v1_auc])
model_metric.add_row(["TF-IDF W2V","Linear kernel", optimal_hyperparameter_tfidfw2v1, tfidfw2v1_auc])
model_metric.add_row(["Bag of Words","RBF kernel", optimal_hyperparameter_bow2, bow_2_auc])
model_metric.add_row(["TF-IDF","RBF kernel", optimal_hyperparameter_tfidf2, tfidf_2_auc])
model_metric.add_row(["Avg W2V","RBF kernel", optimal_hyperparameter_avgw2v2, avgw2v2_auc])
model_metric.add_row(["TF-IDF W2V","RBF kernel", optimal_hyperparameter_tfidfw2v2, tfidfw2v2_auc])
```

```
print(model_metric.get_string(start=0, end=8))
```

Model Name	SVM Type	Hyperparameter	AUC
Bag of Words	Linear kernel	10	0.9386344698096035
TF-IDF	Linear kernel	100	0.9448584206008033
Avg W2V	Linear kernel	100	0.9450313781406859
TF-IDF W2V	Linear kernel	100	0.9450462375470654
Bag of Words	RBF kernel	0.1	0.9131805892863865
TF-IDF	RBF kernel	1	0.9235218931964408
Avg W2V	RBF kernel	1	0.9450313781406859
TF-IDF W2V	RBF kernel	1	0.9235230627445552

9.1 [6.1] Observations:

- 1) Train time: As mentioned in the instructions RBF Kernel has taken significantly higher time and resources to train than the Linear Kernel.
- 2) For Important features for BOW and TF-IDF using Linear kernel is very accurate in terms of the particular words used in the positive and negative review.
- 3) Test AUC scores are very simialr to the train AUC scores for the data which indicates the model is doing a better job on the unseen data