

Keras_Assignment_MNIST

June 5, 2019

```
[83]: import keras
      from keras.datasets import mnist
      from keras.utils import np_utils
      import seaborn as sns
      from keras.initializers import RandomNormal
      from keras.models import Sequential
      from keras.layers import Dense, Activation, Dropout, BatchNormalization

[84]: %matplotlib inline
      import matplotlib.pyplot as plt
      import numpy as np
      import time
      # https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
      # https://stackoverflow.com/a/14434334
      # this function is used to update the plots for each epoch and error
      def plt_dynamic(x, vy, ty, ax, colors=['b']):
          ax.plot(x, vy, 'b', label="Validation Loss")
          ax.plot(x, ty, 'r', label="Train Loss")
          plt.legend()
          plt.grid()
          fig.canvas.draw()

[85]: # https://github.com/keras-team/keras/blob/master/examples/mnist\_mlp.py

      # Importing train and test data from keras
      (X_train, y_train), (X_test, y_test) = mnist.load_data()

[86]: # Reshaping the train and test data from 2D vector to 1D vector
      X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
      X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])

[87]: X_train = X_train.astype('float32')
      X_test = X_test.astype('float32')
      # Min max normalization of train and test data
      X_train = X_train/255
      X_test = X_test/255

[88]: # Converting 10D vector to one-hot encoded features
      Y_train = np_utils.to_categorical(y_train, 10)
```

```
Y_test = np_utils.to_categorical(y_test, 10)
```

```
[89]: # Layer parameters
```

```
output_dim = 10
```

```
input_dim = X_train.shape[1]
```

```
batch_size = 128
```

```
nb_epoch = 20
```

```
[90]: # Defining a model with 2 layers and a dropout layer
```

```
model = Sequential()
```

```
model.add(Dense(512, activation='relu', input_shape=(input_dim,)))
```

```
model.add(Dropout(0.4))
```

```
model.add(Dense(128, activation='relu'))
```

```
model.add(Dense(output_dim, activation='softmax'))
```

```
model.summary()
```

```
-----
Layer (type)                 Output Shape          Param #
-----
dense_140 (Dense)            (None, 512)           401920
-----
dropout_80 (Dropout)         (None, 512)           0
-----
dense_141 (Dense)            (None, 128)           65664
-----
dense_142 (Dense)            (None, 10)            1290
=====
Total params: 468,874
Trainable params: 468,874
Non-trainable params: 0
-----
```

```
[91]: model.compile(optimizer='adam', loss='categorical_crossentropy',
    ↪metrics=['accuracy'])
```

```
history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
    ↪verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 25s 421us/step - loss: 0.2993 -
acc: 0.9092 - val_loss: 0.1187 - val_acc: 0.9635

Epoch 2/20

60000/60000 [=====] - 18s 292us/step - loss: 0.1310 -
acc: 0.9606 - val_loss: 0.0960 - val_acc: 0.9705

Epoch 3/20

60000/60000 [=====] - 18s 302us/step - loss: 0.0965 -

acc: 0.9701 - val_loss: 0.0787 - val_acc: 0.9756
Epoch 4/20
60000/60000 [=====] - 19s 319us/step - loss: 0.0761 -
acc: 0.9764 - val_loss: 0.0804 - val_acc: 0.9737
Epoch 5/20
60000/60000 [=====] - 19s 314us/step - loss: 0.0676 -
acc: 0.9793 - val_loss: 0.0716 - val_acc: 0.9773
Epoch 6/20
60000/60000 [=====] - 19s 325us/step - loss: 0.0584 -
acc: 0.9805 - val_loss: 0.0621 - val_acc: 0.9812
Epoch 7/20
60000/60000 [=====] - 22s 367us/step - loss: 0.0511 -
acc: 0.9830 - val_loss: 0.0624 - val_acc: 0.9821
Epoch 8/20
60000/60000 [=====] - 20s 331us/step - loss: 0.0467 -
acc: 0.9845 - val_loss: 0.0598 - val_acc: 0.9826
Epoch 9/20
60000/60000 [=====] - 20s 337us/step - loss: 0.0425 -
acc: 0.9858 - val_loss: 0.0625 - val_acc: 0.9822
Epoch 10/20
60000/60000 [=====] - 21s 354us/step - loss: 0.0407 -
acc: 0.9863 - val_loss: 0.0659 - val_acc: 0.9802
Epoch 11/20
60000/60000 [=====] - 21s 346us/step - loss: 0.0364 -
acc: 0.9880 - val_loss: 0.0568 - val_acc: 0.9843
Epoch 12/20
60000/60000 [=====] - 21s 349us/step - loss: 0.0343 -
acc: 0.9883 - val_loss: 0.0653 - val_acc: 0.9820
Epoch 13/20
60000/60000 [=====] - 23s 378us/step - loss: 0.0344 -
acc: 0.9887 - val_loss: 0.0630 - val_acc: 0.9826
Epoch 14/20
60000/60000 [=====] - 22s 360us/step - loss: 0.0309 -
acc: 0.9897 - val_loss: 0.0696 - val_acc: 0.9817
Epoch 15/20
60000/60000 [=====] - 23s 377us/step - loss: 0.0298 -
acc: 0.9901 - val_loss: 0.0588 - val_acc: 0.9842
Epoch 16/20
60000/60000 [=====] - 23s 387us/step - loss: 0.0263 -
acc: 0.9907 - val_loss: 0.0618 - val_acc: 0.9849
Epoch 17/20
60000/60000 [=====] - 23s 382us/step - loss: 0.0265 -
acc: 0.9907 - val_loss: 0.0632 - val_acc: 0.9842
Epoch 18/20
60000/60000 [=====] - 23s 383us/step - loss: 0.0254 -
acc: 0.9915 - val_loss: 0.0726 - val_acc: 0.9826
Epoch 19/20
60000/60000 [=====] - 23s 384us/step - loss: 0.0258 -

```
acc: 0.9912 - val_loss: 0.0640 - val_acc: 0.9852
Epoch 20/20
60000/60000 [=====] - 24s 397us/step - loss: 0.0254 -
acc: 0.9919 - val_loss: 0.0631 - val_acc: 0.9839
```

```
[92]: score = model.evaluate(X_test, Y_test, verbose=0)

print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')


# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']

plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06313706721574504

Test accuracy: 0.9839



```
[94]: w_after = model.get_weights()


%matplotlib inline

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



```
[95]: # Defining a model with 3 layers and a dropout layer
model = Sequential()
model.add(Dense(500, activation='relu', input_shape=(input_dim,)))
model.add(Dropout(0.4))
model.add(Dense(300, activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.4))
model.add(Dense(100, activation='relu'))
model.add(BatchNormalization())
model.add(Dense(output_dim, activation='softmax'))
model.summary()
```

```
-----
Layer (type)                 Output Shape          Param #
=====
dense_143 (Dense)            (None, 500)           392500
-----
dropout_81 (Dropout)         (None, 500)           0
-----
dense_144 (Dense)            (None, 300)           150300
-----
batch_normalization_46 (Batc (None, 300)           1200
-----
dropout_82 (Dropout)         (None, 300)           0
-----
dense_145 (Dense)            (None, 100)           30100
-----
batch_normalization_47 (Batc (None, 100)           400
-----
dense_146 (Dense)            (None, 10)            1010
=====
Total params: 575,510
Trainable params: 574,710
Non-trainable params: 800
-----
```

```
[96]: model.compile(optimizer='adam', loss='categorical_crossentropy',
    ↪metrics=['accuracy'])

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
    ↪verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 42s 700us/step - loss: 0.3523 -
```

```

acc: 0.8906 - val_loss: 0.1348 - val_acc: 0.9577
Epoch 2/20
60000/60000 [=====] - 25s 411us/step - loss: 0.1620 -
acc: 0.9496 - val_loss: 0.0975 - val_acc: 0.9694
Epoch 3/20
60000/60000 [=====] - 25s 423us/step - loss: 0.1233 -
acc: 0.9616 - val_loss: 0.0839 - val_acc: 0.9720
Epoch 4/20
60000/60000 [=====] - 25s 415us/step - loss: 0.1024 -
acc: 0.9679 - val_loss: 0.0690 - val_acc: 0.9769
Epoch 5/20
60000/60000 [=====] - 27s 443us/step - loss: 0.0869 -
acc: 0.9725 - val_loss: 0.0717 - val_acc: 0.9777
Epoch 6/20
 512/60000 [...] - ETA: 2:09 - loss: 0.0770 - acc:

/home/mani/anaconda3/lib/python3.7/site-packages/keras/callbacks.py:122:
UserWarning: Method on_batch_end() is slow compared to the batch update
(0.431754). Check your callbacks.
  % delta_t_median)
/home/mani/anaconda3/lib/python3.7/site-packages/keras/callbacks.py:122:
UserWarning: Method on_batch_end() is slow compared to the batch update
(0.216622). Check your callbacks.
  % delta_t_median)

60000/60000 [=====] - 21s 358us/step - loss: 0.0784 -
acc: 0.9750 - val_loss: 0.0642 - val_acc: 0.9806
Epoch 7/20
60000/60000 [=====] - 117s 2ms/step - loss: 0.0727 -
acc: 0.9769 - val_loss: 0.0636 - val_acc: 0.9808
Epoch 8/20
60000/60000 [=====] - 14s 230us/step - loss: 0.0653 -
acc: 0.9789 - val_loss: 0.0598 - val_acc: 0.9828
Epoch 9/20
60000/60000 [=====] - 14s 238us/step - loss: 0.0600 -
acc: 0.9805 - val_loss: 0.0630 - val_acc: 0.9820
Epoch 10/20
1024/60000 [...] - ETA: 1:34 - loss: 0.0568 - acc:

/home/mani/anaconda3/lib/python3.7/site-packages/keras/callbacks.py:122:
UserWarning: Method on_batch_end() is slow compared to the batch update
(0.238180). Check your callbacks.
  % delta_t_median)
/home/mani/anaconda3/lib/python3.7/site-packages/keras/callbacks.py:122:
UserWarning: Method on_batch_end() is slow compared to the batch update
(0.119515). Check your callbacks.
  % delta_t_median)

60000/60000 [=====] - 14s 232us/step - loss: 0.0562 -
acc: 0.9822 - val_loss: 0.0577 - val_acc: 0.9829

```

```

Epoch 11/20
60000/60000 [=====] - 12s 201us/step - loss: 0.0534 -
acc: 0.9826 - val_loss: 0.0572 - val_acc: 0.9823
Epoch 12/20
60000/60000 [=====] - 23s 388us/step - loss: 0.0505 -
acc: 0.9834 - val_loss: 0.0560 - val_acc: 0.9827
Epoch 13/20
60000/60000 [=====] - 24s 395us/step - loss: 0.0460 -
acc: 0.9846 - val_loss: 0.0547 - val_acc: 0.9833
Epoch 14/20
60000/60000 [=====] - 24s 406us/step - loss: 0.0429 -
acc: 0.9858 - val_loss: 0.0516 - val_acc: 0.9845
Epoch 15/20
60000/60000 [=====] - 25s 423us/step - loss: 0.0397 -
acc: 0.9874 - val_loss: 0.0559 - val_acc: 0.9841
Epoch 16/20
60000/60000 [=====] - 19s 310us/step - loss: 0.0395 -
acc: 0.9873 - val_loss: 0.0578 - val_acc: 0.9847
Epoch 17/20
60000/60000 [=====] - 24s 403us/step - loss: 0.0360 -
acc: 0.9885 - val_loss: 0.0614 - val_acc: 0.9832
Epoch 18/20
60000/60000 [=====] - 27s 454us/step - loss: 0.0365 -
acc: 0.9880 - val_loss: 0.0598 - val_acc: 0.9822
Epoch 19/20
60000/60000 [=====] - 29s 487us/step - loss: 0.0348 -
acc: 0.9887 - val_loss: 0.0636 - val_acc: 0.9827
Epoch 20/20
60000/60000 [=====] - 25s 416us/step - loss: 0.0341 -
acc: 0.9889 - val_loss: 0.0572 - val_acc: 0.9843

```

```

[97]: score = model.evaluate(X_test, Y_test, verbose=0)

print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))


vy = history.history['val_loss']
ty = history.history['loss']

plt_dynamic(x, vy, ty, ax)

```


Test score: 0.05719636394771806

Test accuracy: 0.9843



```
[98]: w_after = model.get_weights()

%matplotlib inline


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
```

```
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



```
[99]: # Defining a model with 5 layers and a dropout layer
model = Sequential()
model.add(Dense(500, activation='relu', input_shape=(input_dim,)))
model.add(Dropout(0.3))
model.add(Dense(400, activation='relu'))
model.add(Dropout(0.3))
model.add(Dense(300, activation='relu'))
model.add(Dropout(0.3))
model.add(Dense(200, activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.3))
model.add(Dense(50, activation='relu'))
model.add(BatchNormalization())
model.add(Dense(output_dim, activation='softmax'))
model.summary()
```

Layer (type)	Output Shape	Param #
=====		

dense_147 (Dense)	(None, 500)	392500

dropout_83 (Dropout)	(None, 500)	0

dense_148 (Dense)	(None, 400)	200400

dropout_84 (Dropout)	(None, 400)	0

dense_149 (Dense)	(None, 300)	120300

dropout_85 (Dropout)	(None, 300)	0

dense_150 (Dense)	(None, 200)	60200

batch_normalization_48 (Batch Normalization)	(None, 200)	800

dropout_86 (Dropout)	(None, 200)	0

dense_151 (Dense)	(None, 50)	10050

batch_normalization_49 (Batch Normalization)	(None, 50)	200

dense_152 (Dense)	(None, 10)	510
=====		
Total params: 784,960		
Trainable params: 784,460		
Non-trainable params: 500		

```
[100]: model.compile(optimizer='adam', loss='categorical_crossentropy',
    ↳metrics=['accuracy'])

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
    ↳verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 43s 725us/step - loss: 0.4519 -
acc: 0.8622 - val_loss: 0.1332 - val_acc: 0.9621

Epoch 2/20

60000/60000 [=====] - 33s 542us/step - loss: 0.1658 -
acc: 0.9516 - val_loss: 0.1103 - val_acc: 0.9692

Epoch 3/20

60000/60000 [=====] - 33s 558us/step - loss: 0.1244 -
acc: 0.9624 - val_loss: 0.0918 - val_acc: 0.9739

Epoch 4/20

60000/60000 [=====] - 34s 569us/step - loss: 0.1060 -

acc: 0.9686 - val_loss: 0.0808 - val_acc: 0.9774

Epoch 5/20
60000/60000 [=====] - 36s 599us/step - loss: 0.0885 -
acc: 0.9735 - val_loss: 0.0795 - val_acc: 0.9769

Epoch 6/20
60000/60000 [=====] - 33s 557us/step - loss: 0.0787 -
acc: 0.9768 - val_loss: 0.0880 - val_acc: 0.9751

Epoch 7/20
60000/60000 [=====] - 36s 608us/step - loss: 0.0722 -
acc: 0.9785 - val_loss: 0.0704 - val_acc: 0.9822

Epoch 8/20
60000/60000 [=====] - 37s 621us/step - loss: 0.0618 -
acc: 0.9822 - val_loss: 0.0769 - val_acc: 0.9794

Epoch 9/20
60000/60000 [=====] - 39s 646us/step - loss: 0.0606 -
acc: 0.9820 - val_loss: 0.0735 - val_acc: 0.9808

Epoch 10/20
60000/60000 [=====] - 39s 647us/step - loss: 0.0533 -
acc: 0.9833 - val_loss: 0.0670 - val_acc: 0.9815

Epoch 11/20
60000/60000 [=====] - 39s 656us/step - loss: 0.0484 -
acc: 0.9850 - val_loss: 0.0685 - val_acc: 0.9806

Epoch 12/20
60000/60000 [=====] - 38s 640us/step - loss: 0.0480 -
acc: 0.9857 - val_loss: 0.0659 - val_acc: 0.9816

Epoch 13/20
60000/60000 [=====] - 34s 559us/step - loss: 0.0444 -
acc: 0.9863 - val_loss: 0.0713 - val_acc: 0.9821

Epoch 14/20
60000/60000 [=====] - 40s 661us/step - loss: 0.0391 -
acc: 0.9879 - val_loss: 0.0677 - val_acc: 0.9812

Epoch 15/20
60000/60000 [=====] - 39s 654us/step - loss: 0.0395 -
acc: 0.9882 - val_loss: 0.0665 - val_acc: 0.9828

Epoch 16/20
60000/60000 [=====] - 39s 656us/step - loss: 0.0365 -
acc: 0.9891 - val_loss: 0.0750 - val_acc: 0.9819

Epoch 17/20
60000/60000 [=====] - 40s 663us/step - loss: 0.0351 -
acc: 0.9887 - val_loss: 0.0620 - val_acc: 0.9847

Epoch 18/20
60000/60000 [=====] - 40s 666us/step - loss: 0.0332 -
acc: 0.9899 - val_loss: 0.0645 - val_acc: 0.9839

Epoch 19/20
60000/60000 [=====] - 40s 670us/step - loss: 0.0314 -
acc: 0.9903 - val_loss: 0.0610 - val_acc: 0.9846

Epoch 20/20
60000/60000 [=====] - 41s 689us/step - loss: 0.0298 -

acc: 0.9907 - val_loss: 0.0622 - val_acc: 0.9837

```
[101]: score = model.evaluate(X_test, Y_test, verbose=0)

print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')


# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']

plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06216628487450071

Test accuracy: 0.9837



```
[102]: w_after = model.get_weights()

%matplotlib inline
```

```


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



0.1 Observations

1. The train and validation loss has a very steep difference when only two layers are trained with only Dropout layer.
2. When models trained with multiple Dropout layers and BatchNormalization with 3 and 5 layers the train and validation loss are very close.
3. Weight distributions are smooth gaussian curves in the violin plot for 3 and 5 layered networks.